

# Supervised Learning for Fraud Detection

In this video, we will walk through a comprehensive process of applying supervised machine learning algorithms using real-life data. We will train test and evaluate the following algorithms:

1. Logistic Regression
2. Decision Trees
3. Naive Bayes Classifier

## Import necessary libraries

```
In [1]: # Import necessary libraries
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.model_selection import GridSearchCV
import matplotlib.pyplot as plt
```

## Import the dataset

```
In [2]: # Load data into pandas DataFrame
df = pd.read_csv('C:/Users/Amarkou/Documents/Ecourse/creditcard.csv')
# Select the first 30,000 rows of the DataFrame
df = df.head(60000)

# Print the shape of the sampled DataFrame
print(df)
```

	Time	V1	V2	V3	V4	V5	V6	\
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	
...	...	...	...	...	...	...	...	
59995	49104.0	-1.179821	-0.409590	1.193034	-0.857986	-0.994028	0.013890	
59996	49104.0	1.017824	0.030460	0.592427	1.037257	-0.155280	0.291684	
59997	49104.0	1.019115	-0.332555	1.453262	1.196906	-1.215788	0.327464	
59998	49104.0	1.149634	0.144853	0.349882	0.591884	-0.360690	-0.428566	
59999	49105.0	-2.154248	1.552585	-0.495473	0.143035	-0.736862	0.692048	

	V7	V8	V9	...	V21	V22	V23	\
0	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	
1	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	
2	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412	
3	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	
4	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	
...	...	...	...	...	...	...	...	
59995	1.852891	-0.129020	0.107706	...	0.144079	-0.066047	0.497649	
59996	-0.136690	0.153424	-0.212684	...	-0.088874	-0.381237	0.025885	
59997	-1.061388	0.514961	0.761477	...	0.326159	0.826480	-0.010893	
59998	-0.186583	0.152342	-0.020862	...	-0.210330	-0.688423	0.186908	
59999	-0.959168	1.984428	-0.885507	...	0.217699	0.197859	0.517859	

	V24	V25	V26	V27	V28	Amount	Class
0	0.066928	0.128539	-0.189115	0.133558	-0.021053	149.62	0
1	-0.339846	0.167170	0.125895	-0.008983	0.014724	2.69	0
2	-0.689281	-0.327642	-0.139097	-0.055353	-0.059752	378.66	0
3	-1.175575	0.647376	-0.221929	0.062723	0.061458	123.50	0
4	0.141267	-0.206010	0.502292	0.219422	0.215153	69.99	0
...	...	...	...	...	...	...	...
59995	0.063870	0.354229	0.965341	-0.140004	0.095971	359.40	0
59996	-0.338195	0.231557	-0.582851	0.039972	0.030569	70.35	0
59997	0.160136	0.136173	-0.231982	0.067534	0.029854	27.45	0
59998	0.143570	0.034967	0.098742	-0.020336	0.016581	4.49	0
59999	-0.638766	-0.763836	0.151040	-0.714137	-0.371615	27.63	0

[60000 rows x 31 columns]

## Supervised Learning Modelling Process

## Split data into training and testing sets

```
In [3]: X_train, X_test, y_train, y_test = train_test_split(df.iloc[:, :-1], df.iloc[:]
```

## Define hyperparameters to tune for each algorithm

```
In [4]: lr_params = {'C': [0.1, 1, 10], 'penalty': ['l1', 'l2'], 'solver': ['liblinear']}
dt_params = {'criterion': ['gini', 'entropy'], 'max_depth': [5, 10, 20]}
nb_params = {}
#svm_params = {'C': [0.1, 1, 10], 'kernel': ['linear', 'poly', 'rbf', 'sigmoid']}
#nn_params = {'hidden_layer_sizes': [(50, ), (100, ), (50, 50)], 'activation': [
```

## Train and Test the models

```
In [5]: # Logistic regression
lr = LogisticRegression(max_iter=1000)
lr_gs = GridSearchCV(lr, lr_params, scoring='roc_auc', cv=5)
lr_gs.fit(X_train, y_train)
lr_preds = lr_gs.predict(X_test)

# Decision trees
dt = DecisionTreeClassifier()
dt_gs = GridSearchCV(dt, dt_params, scoring='roc_auc', cv=5)
dt_gs.fit(X_train, y_train)
dt_preds = dt_gs.predict(X_test)

# Naïve Bayes
nb = GaussianNB()
nb_gs = GridSearchCV(nb, nb_params, scoring='roc_auc', cv=5)
nb_gs.fit(X_train, y_train)
nb_preds = nb_gs.predict(X_test)

# SVM
#svm = SVC(probability=True, max_iter=1000)
#svm_gs = GridSearchCV(svm, svm_params, scoring='roc_auc', cv=5)
#svm_gs.fit(X_train, y_train)
#svm_preds = svm_gs.predict(X_test)

# Neural Networks
#nn = MLPClassifier(max_iter=1000)
#nn_gs = GridSearchCV(nn, nn_params, scoring='roc_auc', cv=5)
#nn_gs.fit(X_train, y_train)
#nn_preds = nn_gs.predict(X_test)
```

## Evaluation of algorithms

```
In [6]: print("Logistic Regression Metrics:")
print("Accuracy:", accuracy_score(y_test, lr_preds))
print("Precision:", precision_score(y_test, lr_preds))
print("Recall:", recall_score(y_test, lr_preds))
print("F1 Score:", f1_score(y_test, lr_preds))
print("AUC Score:", roc_auc_score(y_test, lr_preds))
print("\n")

print("Decision Trees Metrics:")
print("Accuracy:", accuracy_score(y_test, dt_preds))
print("Precision:", precision_score(y_test, dt_preds))
print("Recall:", recall_score(y_test, dt_preds))
print("F1 Score:", f1_score(y_test, dt_preds))
print("AUC Score:", roc_auc_score(y_test, dt_preds))
print("\n")

print("Naive Bayes Metrics:")
print("Accuracy:", accuracy_score(y_test, nb_preds))
print("Precision:", precision_score(y_test, nb_preds))
print("Recall:", recall_score(y_test, nb_preds))
print("F1 Score:", f1_score(y_test, nb_preds))
print("AUC Score:", roc_auc_score(y_test, nb_preds))
print("\n")

#print("SVM Metrics:")
#print("Accuracy:", accuracy_score(y_test, svm_preds))
#print("Precision:", precision_score(y_test, svm_preds))
#print("Recall:", recall_score(y_test, svm_preds))
#print("F1 Score:", f1_score(y_test, svm_preds))
#print("AUC Score:", roc_auc_score(y_test, svm.predict_proba(X_test)[:, 1]))
#print("\n")

#print("Neural Networks Metrics:")
#print("Accuracy:", accuracy_score(y_test, nn_preds))
#print("Precision:", precision_score(y_test, nn_preds))
#print("Recall:", recall_score(y_test, nn_preds))
#print("F1 Score:", f1_score(y_test, nn_preds))
#print("AUC Score:", roc_auc_score(y_test, nn.predict_proba(X_test)[:, 1]))
#print("\n")
```

Logistic Regression Metrics:  
Accuracy: 0.9981666666666666  
Precision: 0.8611111111111112  
Recall: 0.5254237288135594  
F1 Score: 0.6526315789473686  
AUC Score: 0.7625725187738718


Decision Trees Metrics:  
Accuracy: 0.9990555555555556  
Precision: 0.9565217391304348  
Recall: 0.7457627118644068  
F1 Score: 0.8380952380952381  
AUC Score: 0.8728256176790402

Naive Bayes Metrics:  
Accuracy: 0.985  
Precision: 0.1607717041800643  
Recall: 0.847457627118644  
F1 Score: 0.2702702702702703  
AUC Score: 0.9164549715215315

## ROC Curve graph

```
In [7]: # Evaluation of algorithms (continued)
fpr_lr, tpr_lr, thresholds_lr = roc_curve(y_test, lr_preds)
fpr_dt, tpr_dt, thresholds_dt = roc_curve(y_test, dt_preds)
#fpr_svm, tpr_svm, thresholds_svm = roc_curve(y_test, svm.predict_proba(X_test)[:,1])
fpr_nb, tpr_nb, thresholds_nb = roc_curve(y_test, nb_preds)
#fpr_nn, tpr_nn, thresholds_nn = roc_curve(y_test, nn.predict_proba(X_test)[0,:])

plt.figure(figsize=(8, 6))
plt.plot(fpr_lr, tpr_lr, label='Logistic Regression')
plt.plot(fpr_dt, tpr_dt, label='Decision Trees')
#plt.plot(fpr_svm, tpr_svm, label='SVM')
plt.plot(fpr_nb, tpr_nb, label='Naïve Bayes')
#plt.plot(fpr_nn, tpr_nn, label='Neural Networks')
plt.plot([0, 1], [0, 1], color='black', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()
```



In [ ]: